

# Foundations of Machine Learning and AI

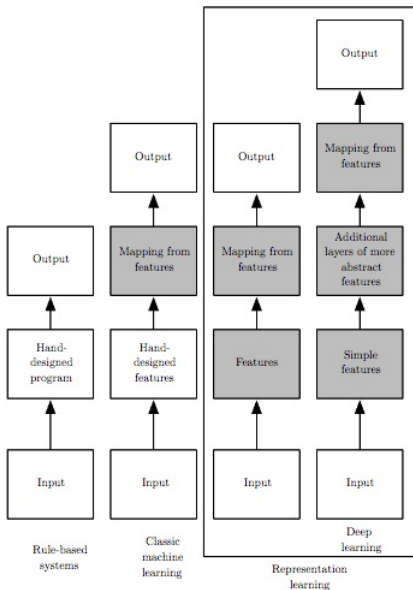
Theodoros Evgeniou - Nicolas Vayatis

Sessions 5-6: Data representations, feature learning and applications

## What we have seen so far

- Machine Learning is about learning (= choosing = estimating) a function from data
- The key concept is the complexity of the function space ("hypothesis space") where we look for our solution ("how many functions we select from")
- The art of learning is to use the data to adjust the complexity of the hypothesis space - while implicitly considering the *approximation error*.
- In the particular case of least square linear regression, complexity calibration can (also) be achieved by only selecting and using a small subset of the variables (the problem of variable selection).

## Another "Big picture" of Learning



## Objectives for this class

- Focus on **feature selection** and **feature learning**: learning ("finding" or "choosing") a representation of the data
- Develop new regularisation/machine learning formulations for other applications such as learning (= estimating the missing entries of) matrices - for example used in recommender systems
- We will also learn about some **optimization** approaches to solve machine learning formulations/methods (possibly nonconvex optimization problems): **Optimization is central for machine learning**

## What we know about sparsity from sessions 3-4

- Sparsity-inducing methods: LASSO
- Motivation in linear predictive models: relaxation of  $\ell_0$  constraint on number of independent variables used, namely from minimizing

$$\|\mathbf{Y} - \mathbf{X}\beta\|^2 + \lambda\|\beta\|_0$$

to minimizing

$$\|\mathbf{Y} - \mathbf{X}\beta\|^2 + \lambda\|\beta\|_1$$

- Advantages: tractable computations, interpretable models
- Byproduct: sparsistency (i.e. how many, and which variables to use)

Application (today): Matrix completion  
with (rank) Sparsity  
("Netflix Recommendation Competition")

		Item			
		W	X	Y	Z
User	A		4.5	2.0	
	B	4.0		3.5	
	C		5.0		2.0
	D		3.5	4.0	1.0

Rating Matrix

# Feature Selection and Learning

- A. Feature "Learning": PCA and variants
- B. Feature Selection: LASSO with optimization methods
- C. Applications: sparse coding, (kernels), matrix completion

## A. Feature "Learning": Principal Component Analysis (PCA) and variants

### 1. Classical PCA




# What all students should know

## PCA

- Motivation: Dimensionality reduction
- Principle: Find an orthogonal basis to represent (project on) the data, which captures the directions of highest dispersion (variance) of the data
- Underlying assumption: Gaussian, highly correlated data

# Idea of PCA



# PCA



## Classical construction

- Compute the covariance (or correlation) matrix of the data
- Find the eigen-elements (values/vectors) - eigenvectors being orthogonal - of this matrix
- Principal components are ordered from the larger eigenvalue to the smallest
- Dimensionality reduction from  $d$  to (small)  $r$  is performed by projecting the initial data points on the first (principal)  $r$  eigenvectors

## PCA applied to music recommendation

LastFM data set <http://www.last.fm>:


[//anthony.liekens.net/index.php/Computers/DataMining](http://anthony.liekens.net/index.php/Computers/DataMining)




# PCA applied to time series

## Job hiring data (1/3)

JOLTS data set available at <https://www.bls.gov/jlt/>






Principal Components (normalized data)



# PCA applied to time series

## Job hiring data (2/3)




### Components interpretation



# PCA applied to time series

## Job hiring data (3/3)

### Projection on principal components



# PCA applied to time series

## Financial data (1/2)

Paper by Avellenada and Lee (2008)




Figure 1: Eigenvalues of the correlation matrix of market returns computed on May 1 2007 estimated using a 1-year window (measured as percentage of explained variance)



## PCA applied to time series

### Financial data (2/2)

Paper by Avellenada and Lee (2008)




Figure 4: First eigenvector sorted by coefficient size. The x-axis shows the ETF corresponding to the industry sector of each stock.

## A different view on PCA

- Denote by  $X$  the data matrix of size  $d \times n$  (assume that the points are centered) and by  $\|M\|_F^2 = \sum_{i,j} M_{ij}^2$  the square of the *Frobenius norm* of the matrix  $M = (M_{ij})_{ij}$
- Solve the minimization problem:

$$\min_{P,Z} \|X - PZ\|_F^2 \text{ subject to } P^T P = I_r$$

where  $P$  is the projection matrix of size  $d \times r$  (the matrix whose columns are the first  $r$  eigenvectors), and  $Z$  is  $r \times n$  matrix of the projected points in the  $r$ -dimensional subspace. We also have the *orthogonality* constraint  $P^T P = I_r$  (eigenvectors are orthogonal)

## A low-rank formulation of PCA

- An alternative formulation to the previous optimization problem, by setting:  $A = PZ$ , is:

$$\min_A \|X - A\|_F^2 \text{ subject to } \text{rank}(A) = r$$

- Theoretical result (Vidal, Ma, Sastry (2016)): an optimal solution to this problem is given by:

$$A = U_r \Sigma_r V_r$$

where  $U_r$  and  $V_r$  have orthogonal columns of size  $d \times r$  and  $n \times r$  respectively,  $\Sigma_r$  diagonal square matrix of size  $r \times r$ . The matrices  $U_r$ ,  $\Sigma_r$ ,  $V_r$  correspond to the **reduced singular value decomposition (SVD) of matrix  $X$** .

## Some linear algebra background: SVD decomposition

A generalization of eigenvalues and eigenvectors.

- Definition:  $\sigma$  is a singular value of a rectangular  $d \times n$  matrix  $X$  if there exist unit two vectors  $u \in \mathbb{R}^d$  and  $v \in \mathbb{R}^n$  such that

$$X^T u = \sigma v \quad \text{and} \quad Xv = \sigma u$$

The vectors  $u$  and  $v$  are called **singular vectors**.

- Theorem: For any rectangular matrix, there exist  $U$  and  $V$  orthogonal matrices of size  $d \times d$  and  $n \times n$  respectively and a diagonal matrix  $\Sigma$  of size  $d \times n$  such that:

$$X = U\Sigma V^T$$

A. Feature "Learning":

Principal Component Analysis (PCA) and variants

2. Nonnegative Matrix Factorization (NMF)

## Some issues with PCA

- PCA is sensitive to outliers; empirical covariance matrix converges to real covariance slowly wrt sample size...
- What if natural components are not Gaussian? what if they are not orthogonal but independent (check more than just their correlation)? ...
- What about interpretation? Maybe we need nonnegativity of matrix  $Z$  (the new data representation) → Nonnegative Matrix Factorization

# Nonnegative Matrix Factorization

PCA




$\times$



$\Pi$



NMF



$\times$



$=$

Original



D.D. Lee and H. S.Seung, "Learning the parts of objects by non-negative matrix factorization", Nature 401 (6755), pp. 788–791, 1999

A. Feature "Learning":  
Principal Component Analysis (PCA) and variants

3. Robust PCA



# A Machine Learning-type Formulation

## Robust PCA

- Introduced by Candès, Li, Ma, Wright (2011)
- Motivation: assume a decomposition of the data matrix  $X = L + S$  where  $L$  is low rank and  $S$  is sparse.
- *Principal Component Pursuit*: the *nuclear norm* (also called *Trace norm*)  $\|\cdot\|_*$  defined as the sum of singular values; note with  $\|\cdot\|_1$  the  $\ell_1$  matrix norm (sum of the absolute values of all the entries of the matrix). We search for matrices  $L$  and  $S$ :

$$\min_{L, S} \|L\|_* + \lambda \|S\|_1 \text{ subject to } L + S = X$$

- Main theoretical result: under some assumptions the *exact* solution may be recovered by this procedure

## Other variants of PCA

- Sparse PCA
- Nonlinear PCA, Kernel PCA
- Multidimensional scaling, Local embeddings, Laplacian eigenmaps
- ...

Reference: book by Vidal, Ma, Sastry. Generalized Principal Component Analysis. Springer (2016)

## B. Feature selection: LASSO with optimization methods

# The LASSO for linear models

## From $\ell_0$ to $\ell_1$

- Consider the LASSO estimation (learning) method: for any  $\lambda > 0$ ,


$$\hat{\beta}_\lambda \in \arg \min_{\beta \in \mathbb{R}^d} \{ \|\mathbf{Y} - \mathbf{X}\beta\|^2 + \lambda \|\beta\|_1 \}$$

where the  $\ell_1$ -norm is:

$$\|\beta\|_1 = \sum_{j=1}^d |\beta_j|$$

# Blessings of the LASSO

- Approximate solutions via efficient algorithms building the so-called *regularization path* (find for all values of  $\lambda$  the  $\hat{\beta}(\lambda)$ ):



- Theoretical soundness: it can be shown that (if the real model is linear): as  $n, d \rightarrow \infty$

$$\frac{1}{n} \mathbb{E}(\|\mathbf{x}\beta^* - \mathbf{x}\hat{\beta}\|^2) \leq C \|\beta^*\|_1 \sqrt{\frac{\log d}{n}}$$

# Optimization methods for LASSO estimation

[mainly pointers to different approaches and literature]

- Least Angle Regression
- Coordinate Descent
- Proximal methods

# Optional material

Optimization methods applied to LASSO

## First algorithm: Least Angle Regression (LARS)

- LARS = variant of the incremental stagewise procedure for adding variables in a linear model
  - Least Angle Regression paper by Efron-Hastie-Johnstone-Tibshirani (AoS, 2004)
  - Previous work by Osborne et al. (2000) on the so-called homotopy method
  - Also related to greedy approaches such as Orthogonal Matching Pursuit (by Mallat, Zhang (1993), Mallat, Davis, Zhang (1994))
- Recovers the full regularization path  $\lambda \rightarrow \hat{\beta}(\lambda)$  of the LASSO
- Success of the procedure based on the fact that LASSO path is piecewise linear.
- Computational efficiency: one ordinary least square computation at each step




## Least Angle Regression: Pseudocode

- 1 Start with all coefficients  $\beta$  equal to zero.
- 2 Find the predictor  $x_j$  most correlated with  $y$
- 3 Increase the coefficient  $\beta_j$  in the direction of the sign of its correlation with  $y$  until some other predictor  $x_k$  has as much correlation with  $r = y - \hat{y}$  as  $x_j$  has.
- 4 Increase  $(\beta_j, \beta_k)$  in their joint least squares direction, until some other predictor  $x_m$  has as much correlation with the residual  $r$ .
- 5 Continue until: all predictors are in the model (corresponding to the solution when  $\lambda$  is small)

## Second algorithm: Coordinate Descent

- Simple idea of one dimensional optimization with cyclic iteration over all variables, until convergence
- Optimization at each step amounts to a one-dimensional LASSO problem
- Solution obtained as a soft thresholding of the one-dimensional ordinary least square estimate.



## Third algorithm: Proximal methods

- Parikh-Boyd tutorial paper (2013): "Much like Newton's method is a standard tool for solving unconstrained smooth optimization problems of modest size, proximal algorithms can be viewed as an analogous tool for nonsmooth, constrained, large-scale, or distributed versions of these problems."
- Early work goes back to Moreau (1960s) then Nemirovski, Yudin (1983)
- Rediscovered around 2005 with applications to signal processing and solving certain optimization problems

# Proximal method

## Principle

- Applies to a problem of the form:

$$\min_{\beta} \{L(\beta) + \psi(\beta)\}$$

when:  $L$  is smooth, convex, with "bounded" gradient, and  $\psi$  is continuous, convex, but non-smooth

- The proximal algorithm is a descent algorithm which provides a sequence  $\beta_t$  obtained as follows: at each step  $t$ ,

$$\beta_t = \text{prox}(\psi, \beta_{t-1} - \nabla L(\beta_{t-1}))$$

where  $\text{prox}$  is the so-called *proximal operator* (generalizes the concept of orthogonal projection)

# Proximal method (1/2)

## Definition of proximal operator

- Definition of the proximal operator for the nonsmooth term  $\psi$  of the objective  $L + \psi$

$$\text{prox}(\psi, z) = \arg \min_{\beta} \left\{ \frac{1}{2} \|\beta - z\|_2^2 + \psi(\beta) \right\}$$

- Interpretation: The proximal operator finds a point that corresponds to a trade-off between minimizing  $\psi$  and being near to the point  $z$ .

## Proximal method (2/2)

### Application to LASSO

- Here:  $L(\beta) = \frac{1}{2}\|X\beta - y\|_2^2$  and  $\psi(\beta) = \lambda\|\beta\|_1$
- Gradient step relies on the gradient of the smooth term  $L$ :

$$\nabla L(\beta) = X^T(X\beta - y)$$

- Proximal operator for the  $\ell_1$  norm is given by:

$$\text{prox}(\lambda\|\cdot\|_1, z) = (z - \lambda)_+ - (-z - \lambda)_+$$

(soft thresholding operator on each component of  $z$ )

- Also called ISTA (for Iterative Shrinkage Thresholding Algorithm)

End of optional material

## C. Applications:

### 1. Sparse coding



## Motivations and references

- Some features (to represent the data) may be good for compression but not for interpretation (and vice versa); they may also simply fail to "lead to" sparse representations (e.g., learn functions that use only a few of the features)
- Can we learn data features (representation) so that the functions we learn (estimate) in that representation ("space") are also sparse?
- Idea is to exploit the fact that *similar patterns may be repeated in the data (even if they are not smooth)*
- (Can also be used to handle some cases of non-stationarity)

References: Olshausen and Field (1997) Kreutz-Delgado et al. (2003), Mairal, Elad, Sapiro (2008), Gribonval et al. (2015)

# Sparse convolutional coding

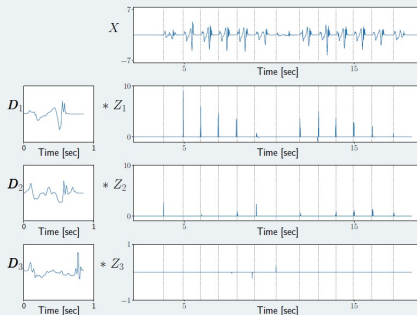
## Notation

- ▶  $X$  is a signal of length  $T$
- ▶  $\mathcal{E}$  is a noise signal of length  $T$
- ▶  $D$  is a set of  $K$  patterns of length  $W$
- ▶  $Z$  is a signal of length  $L = T - W + 1$  in  $\mathbb{R}^K$

## Sparse Convolutional model:

$$X[t] = \sum_{k=1}^K (D_k * Z_k)[t] + \mathcal{E}[t]$$

with  $Z$  sparse. Few of its coefficients are non-zero.



# Sparse (linear) coding

## Notation


- ▶  $x$  is a vector in  $\mathbb{R}^T$
- ▶  $\epsilon$  is a noise vector in  $\mathbb{R}^T$
- ▶  $D$  is a matrix in  $\mathbb{R}^{T \times LK}$
- ▶  $z$  is a coding vector in  $\mathbb{R}^{LK}$

## Sparse Linear model:

$$x = Dz + \epsilon$$

with  $z$  sparse. Few of its coefficients are non-zero.

## Link with convolutional model



# Sparse coding Formulation

- Objective: find both the features  $D$  and the activations  $Z$  that yield to the sparse representation of the data  $X$  up to some error  $\varepsilon$
- Formulation:

$$\min_{D,Z} \left\{ \sum_{i=1}^n \|Z_i\|_0 \right\} \text{ subject to } \|X - DZ\|_2 \leq \varepsilon$$


## Sparse coding

### Towards nonconvex optimization

- Same complexity as  $\ell_0$  norm minimization problem. In practice, it is solved with an  $\ell_1$ -type relaxation
- But: for fixed  $D$ , minimization over  $Z$  is convex while the joint optimization wrt both  $Z$  and  $Z$  is not convex
- Main strategy for non convex matrix factorization problems: alternating minimization (Douglas-Rachford) or Block coordinate descent

## Sparse coding: Examples

- Images (text? multimedia?, etc)



- Representation of consumer products ("meta-attributes") and utility functions (see also conjoint analysis and Multi-task Learning in Sessions 13-14).

## C. Applications

### 2. Matrix completion

# Matrix completion: Recommender Systems Application

	Item			
	W	X	Y	Z
A		4.5	2.0	
B	4.0		3.5	
C		5.0		2.0
D		3.5	4.0	1.0

Rating Matrix

=

A	1.2	0.8
B	1.4	0.9
C	1.5	1.0
D	1.2	0.8

User Matrix

X

	W	X	Y	Z
	1.5	1.2	1.0	0.8
	1.7	0.6	1.1	0.4

Item Matrix



## Matrix completion: Problem statement

- Original optimization formulation (kind of "Ivanov Regularization" with no error on the available matrix entries - our data)

$$\min_X \{\text{rank}(X)\} \text{ subject to } X_{ij} = M_{ij}, \forall (i,j) \in \Omega$$

where  $\Omega = \{(i,j) : M_{ij} \text{ the available data}\}$ .

- **Key Challenge:** Non-convex problem, hard to solve

## Matrix completion: Convex Relaxation

- Recall the *nuclear norm* of  $X$  is  $\|X\|_* = \sum_{i=1}^{\min(n,m)} \sigma_i$ , where  $\sigma_i$  are the singular values of  $X$  (recall the SVD of  $X$  is  $X = U\Sigma V^T$ )
- Convex formulation of the matrix completion problem:

$$\min_X \|X\|_* \text{ subject to } X_{ij} = M_{ij}, \forall (i,j) \in \Omega$$

where  $\Omega = \{(i,j) : M_{ij} \text{ the available data}\}$ .

- Regularization formulation:** Nuclear norm penalty

$$\min_X \left\{ \frac{1}{2} \sum_{ij \in \Omega} (X_{ij} - M_{ij})^2 + \lambda \|X\|_* \right\}$$

# Optional material

Resolution of the matrix completion problem

## Matrix completion Solution (1/2)

- Simplified problem (no mask  $\Omega$ ):

$$\min_X \left\{ \frac{1}{2} \|X - M\|^2 + \lambda \|X\|_* \right\}$$

- The solution is closed form and given by:

$$\text{shrink}(X, \lambda) = U \Sigma(\lambda) V^T$$

where  $\Sigma(\lambda) = \text{diag}((\sigma_i - \lambda)_+)$

- Note: the solution uses only the singular values that are larger than  $\lambda$ ...

## Matrix completion Solution (2/2)


- Need a trick to deal with the  $\Omega$
- Use an auxiliary matrix  $Y$  which is complete
- Define  $\Pi_{\Omega}(X)$  the matrix with coefficients  $X_{ij}$  if  $(i,j) \in \Omega$  and zero if  $(i,j) \notin \Omega$
- Iterative algorithm (called "SVT"):
  - ① Set  $\lambda > 0$  and sequence of step sizes  $(\delta_k)_{k \geq 1}$
  - ② Start with  $Y_0 = 0$  matrix of size  $n \times m$
  - ③ At each step  $k$ , compute:

$$\begin{cases} X_k &= \text{shrink}(Y_{k-1}, \lambda) \\ Y_k &= Y_{k-1} + \delta_k \Pi_{\Omega}(M - X_k) \end{cases}$$

End of optional material

Coming next

## Another "Big picture" of Learning





## Next sessions

- Deep Learning (also learn "hierarchical representations" of the data - what if "nature" has this structure?)
- More on optimization (e.g., stochastic gradient)
- Q&A, catch up, class projects
- Overview of popular algorithms (sessions 9-10)